

Llama 3.1 8B (Dense) vs DeepSeek-v2 Lite (MoE) — Decode Throughput (tokens/s)

